

Zhengpu Zhao CS285 HW1 Report

TLDR:

Submitting the PDF. Make a PDF report containing:
Table 1 for a table of results from Question 1.2,
Figure 1 for Question 1.3. and
Figure 2 with results from question 2.2.

Table 1 for a table of results from Question 1.2:

| | bc_ant | expert_ant | bc_walker2d | expert_walker2d |
|--------------------|------------------|------------------|-------------------|------------------|
| Average Return | 4669.265625 | 4713.6533203125 | 289.3726501464844 | 5566.845703125 |
| Standard Deviation | 81.988525390625 | 12.196533203125 | 86.68399810791016 | 9.237548828125 |
| Min Return | 4585.6728515625 | 4701.45654296875 | 6.183746337890625 | 5557.6083984375 |
| Max Return | 4798.67333984375 | 4725.849609375 | 440.492919921875 | 5576.08349609375 |

BC_Ant Achievement as percentage of Expert_Ant: 99%
BC_Walker2d Achievement as percentage of Expert_Walker2d: 5.2%

Figure 1 for Question 1.3:

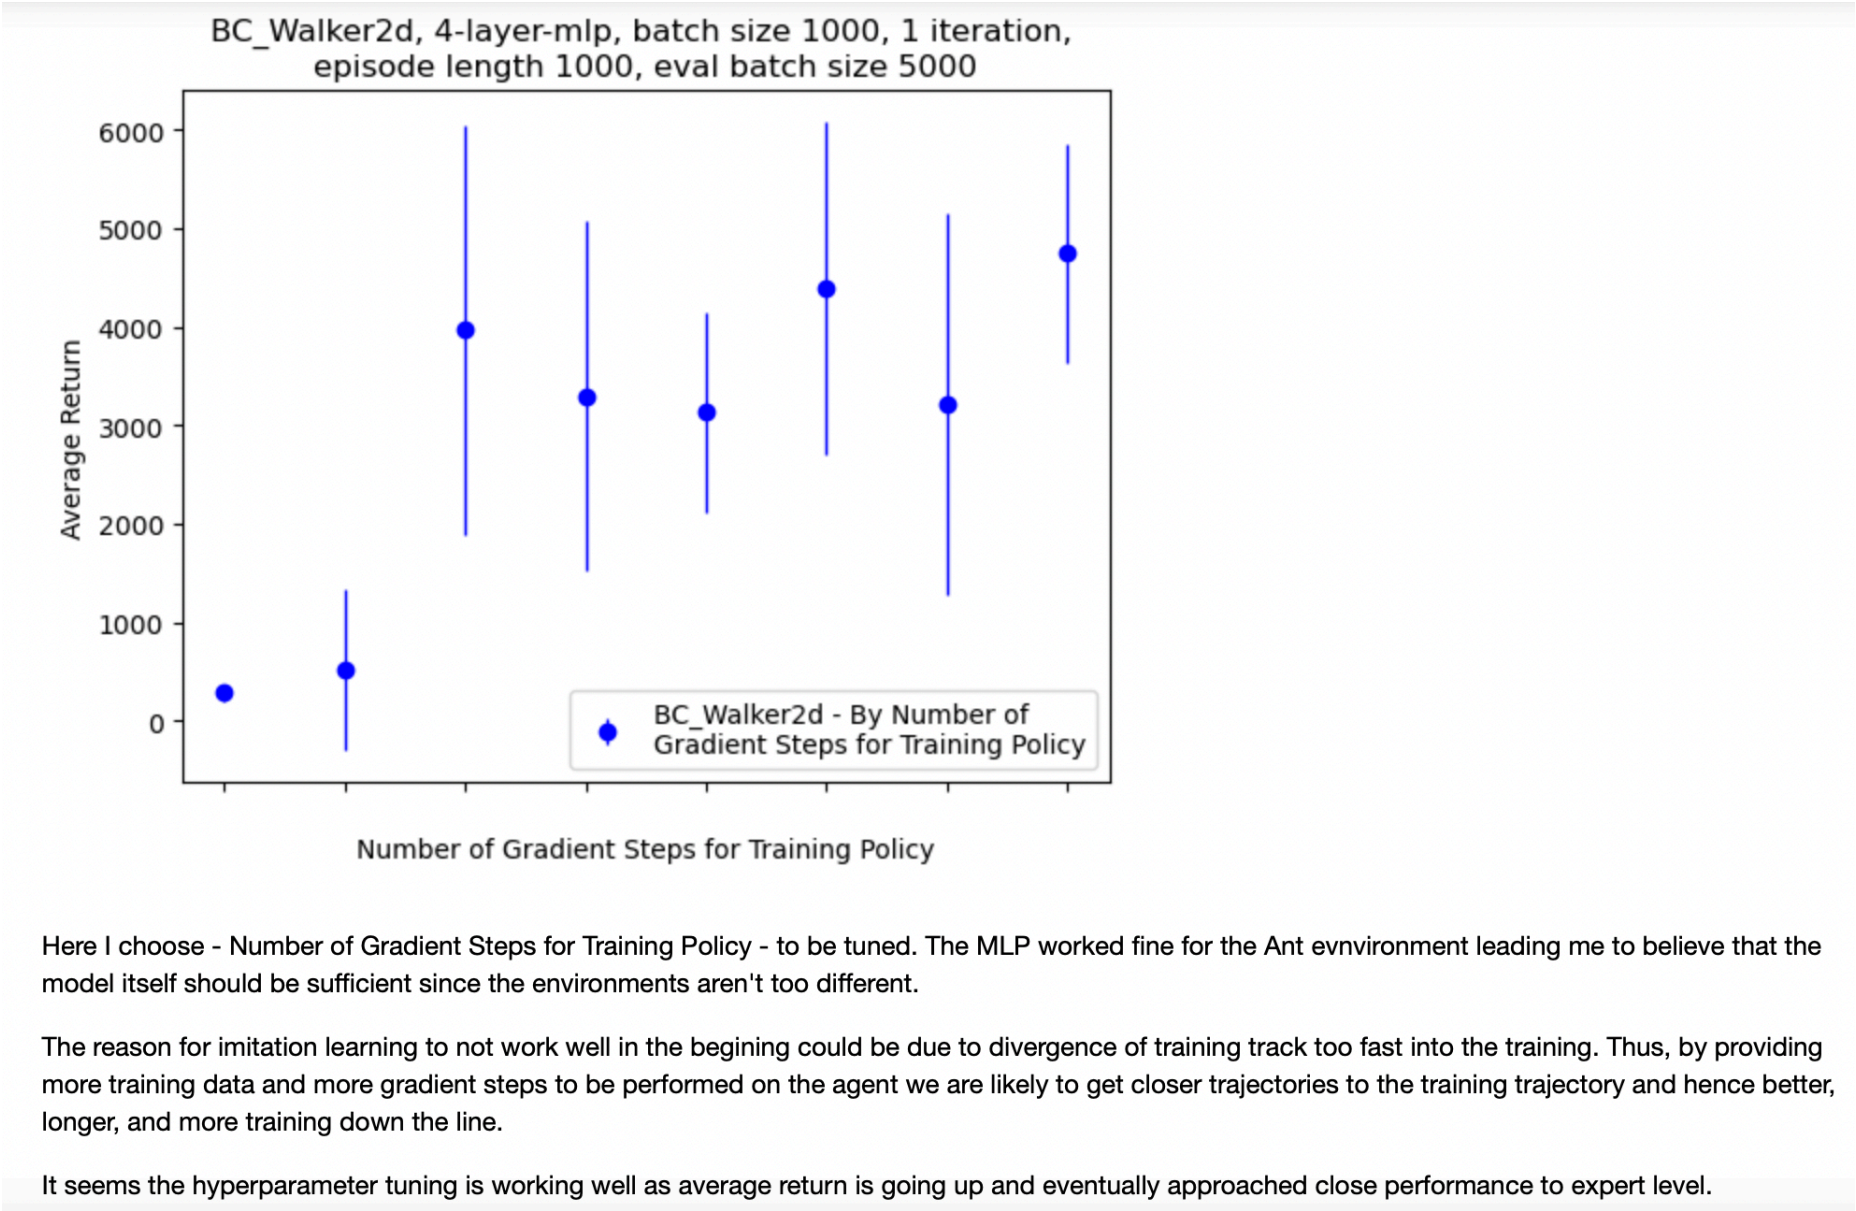



Figure 2 with results from question 2.2:



```
In [1]: 1 import os
2 import numpy as np
3 import pandas as pd
4
5 from collections import defaultdict
6 from tensorboard.backend.event_processing.event_accumulator import EventAccumulator
```

```
In [2]: 1 def tabulate_events(dpath):
2     summary_iterators = [EventAccumulator(os.path.join(dpath, dname)).Reload() for dname in os.listdir(dpath)]
3
4     tags = summary_iterators[0].Tags()['scalars']
5
6     for it in summary_iterators:
7         assert it.Tags()['scalars'] == tags
8
9     out = defaultdict(list)
10    steps = []
11
12    for tag in tags:
13        steps = [e.step for e in summary_iterators[0].Scalars(tag)]
14
15        for events in zip(*[acc.Scalars(tag) for acc in summary_iterators]):
16            assert len(set(e.step for e in events)) == 1
17
18            out[tag].append([e.value for e in events])
19
20    return out, steps
21
22
23 def to_df(dpath):
24     dirs = os.listdir(dpath)
25
26     d, steps = tabulate_events(dpath)
27     tags, values = zip(*d.items())
28     np_values = np.array(values)
29
30     whole_df = pd.DataFrame()
31
32     for index, tag in enumerate(tags):
33         values_col = np_values[index]
34         sq_val = values_col.squeeze()
35         if sq_val.shape == ():
36             a_col = sq_val.reshape((1,))
37         else:
38             a_col = sq_val
39         whole_df[tag] = a_col
40
41     return whole_df
42
43 def get_file_path(dpath, tag):
44     file_name = tag.replace("/", "_") + '.csv'
45     folder_path = os.path.join(dpath, 'csv')
46     if not os.path.exists(folder_path):
47         os.makedirs(folder_path)
48     return os.path.join(folder_path, file_name)
```

Use:

path = "Path"

to_csv(path)

Q1.2

```
python cs285/scripts/run_hw1.py
--expert_policy_file cs285/policies/experts/Ant.pkl
--env_name Ant-v4 --exp_name bc_ant --n_iter 1
--expert_data cs285/expert_data/expert_data_Ant-v4.pkl
--video_log_freq -1 --eval_batch_size 5000
--ep_len 1000

python cs285/scripts/run_hw1.py
--expert_policy_file cs285/policies/experts/Walker2d.pkl
--env_name Walker2d-v4 --exp_name bc_walker2d --n_iter 1
--expert_data cs285/expert_data/expert_data_Walker2d-v4.pkl
--video_log_freq -1 --eval_batch_size 5000
--ep_len 1000
```

BC_Ant

q1_bc_ant_Ant-v4_12-09-2022_00-35-43

```
In [3]: 1 bc_ant_exp_path = "q1_bc_ant_Ant-v4_12-09-2022_00-35-43"
        2 bc_ant_exp_df = to_df(bc_ant_exp_path)
        3 # ignore error

-----
ImportError                                Traceback (most recent call last)
File /opt/anaconda3/envs/cs285/lib/python3.8/site-packages/tensorboard/compat/__init__.py:42, in tf()
      41 try:
----> 42     from tensorboard.compat import notf # noqa: F401
      43 except ImportError:

ImportError: cannot import name 'notf' from 'tensorboard.compat' (/opt/anaconda3/envs/cs285/lib/python3.8/site-packages/tensorboard/compat/__init__.py)

During handling of the above exception, another exception occurred:

RuntimeError                                Traceback (most recent call last)
RuntimeError: module compiled against API version 0xf but this version of numpy is 0xe
```

```
In [4]: 1 bc_ant_exp_df.head()
```

Out[4]:

| | Eval_AverageReturn | Eval_StdReturn | Eval_MaxReturn | Eval_MinReturn | Eval_AverageEpLen | Train_AverageReturn | Train_StdReturn | Train_MaxReturn |
|---|--------------------|----------------|----------------|----------------|-------------------|---------------------|-----------------|-----------------|
| 0 | 4669.265625 | 81.988525 | 4798.67334 | 4585.672852 | 1000.0 | 4713.65332 | 12.196533 | 4725 |

```
In [5]: 1 bc_ant_exp_Eval_AverageReturn = bc_ant_exp_df["Eval_AverageReturn"][0]
        2 bc_ant_exp_Train_AverageReturn = bc_ant_exp_df["Train_AverageReturn"][0]
        3 achieve_percentage = bc_ant_exp_Eval_AverageReturn / bc_ant_exp_Train_AverageReturn
        4 print("{} %".format(achieve_percentage*100))

99.05831650534797 %
```

BC_Walker2d

q1_bc_walker2d_Walker2d-v4_12-09-2022_00-36-44

```
In [6]: 1 bc_walker2d_exp_path = "q1_bc_walker2d_Walker2d-v4_12-09-2022_00-36-44"
        2 bc_walker2d_exp_df = to_df(bc_walker2d_exp_path)
        3 bc_walker2d_exp_df.head()
```

Out[6]:

| | Eval_AverageReturn | Eval_StdReturn | Eval_MaxReturn | Eval_MinReturn | Eval_AverageEpLen | Train_AverageReturn | Train_StdReturn | Train_MaxReturn |
|---|--------------------|----------------|----------------|----------------|-------------------|---------------------|-----------------|-----------------|
| 0 | 289.37265 | 86.683998 | 440.49292 | 6.183746 | 126.650002 | 5566.845703 | 9.237549 | 5576 |

```
In [7]: 1 bc_walker2d_exp_Eval_AverageReturn = bc_walker2d_exp_df["Eval_AverageReturn"][0]
        2 bc_walker2d_exp_Train_AverageReturn = bc_walker2d_exp_df["Train_AverageReturn"][0]
        3 achieve_percentage = bc_walker2d_exp_Eval_AverageReturn / bc_walker2d_exp_Train_AverageReturn
        4 print("{} %".format(achieve_percentage*100))

5.198143896534484 %
```

```
In [8]:
```

```

In [8]: from tabulate import tabulate
import matplotlib.pyplot as plt

3
4
5 means = ['Average Return',
6         bc_ant_exp_Eval_AverageReturn,
7         bc_ant_exp_Train_AverageReturn,
8         bc_walker2d_exp_Eval_AverageReturn,
9         bc_walker2d_exp_Train_AverageReturn]
10
11 std = ['Standard Deviation',
12        bc_ant_exp_df["Eval_StdReturn"][0],
13        bc_ant_exp_df["Train_StdReturn"][0],
14        bc_walker2d_exp_df["Eval_StdReturn"][0],
15        bc_walker2d_exp_df["Train_StdReturn"][0]]
16
17 mins = ['Min Return',
18        bc_ant_exp_df["Eval_MinReturn"][0],
19        bc_ant_exp_df["Train_MinReturn"][0],
20        bc_walker2d_exp_df["Eval_MinReturn"][0],
21        bc_walker2d_exp_df["Train_MinReturn"][0]]
22
23 maxes = ['Max Return',
24         bc_ant_exp_df["Eval_MaxReturn"][0],
25         bc_ant_exp_df["Train_MaxReturn"][0],
26         bc_walker2d_exp_df["Eval_MaxReturn"][0],
27         bc_walker2d_exp_df["Train_MaxReturn"][0]]
28
29
30 # create stacked errorbars:
31 plt.errorbar(np.array(1), np.array(means[1]), np.array(std[1]), color='blue', fmt='ok', lw=1)
32 plt.errorbar(np.array(2), np.array(means[2]), np.array(std[2]), color='green', fmt='ok', lw=1)
33 plt.errorbar(np.array(3), np.array(means[3]), np.array(std[3]), color='red', fmt='ok', lw=1)
34 plt.errorbar(np.array(4), np.array(means[4]), np.array(std[4]), color='orange', fmt='ok', lw=1)
35 plt.legend(['bc_ant (99% of expert)', 'expert_ant',
36            'nbc_walker2d (5.2% of expert)', 'nexpert_walker2d'])
37 plt.xlabel("Agent")
38 plt.ylabel("Return")
39 plt.xticks(color='w')
40 plt.title("4-layer-mlp, batch size 1000, 1 iteration of behavior cloning, \nepisode length 1000, eval bat
41 plt.show()
42
43 Q1_2_table = [[' ', 'bc_ant', 'expert_ant', 'bc_walker2d', 'expert_walker2d'],
44              means, std, mins, maxes]
45 print(tabulate(Q1_2_table, tablefmt='fancy_grid'))
46
47 print('BC_Ant Achievement as percentage of Expert_Ant: 99%')
48 print('BC_Walker2d Achievement as percentage of Expert_Walker2d: 5.2%')

```

/var/folders/80/mq_tqq2929b94txcm70_rbyr0000gn/T/ipykernel_93228/4274029093.py:31: UserWarning: color is redundantly defined by the 'color' keyword argument and the fmt string "ok" (-> color='k'). The keyword argument will take precedence.

```
plt.errorbar(np.array(1), np.array(means[1]), np.array(std[1]), color='blue', fmt='ok', lw=1)
```

/var/folders/80/mq_tqq2929b94txcm70_rbyr0000gn/T/ipykernel_93228/4274029093.py:32: UserWarning: color is redundantly defined by the 'color' keyword argument and the fmt string "ok" (-> color='k'). The keyword argument will take precedence.

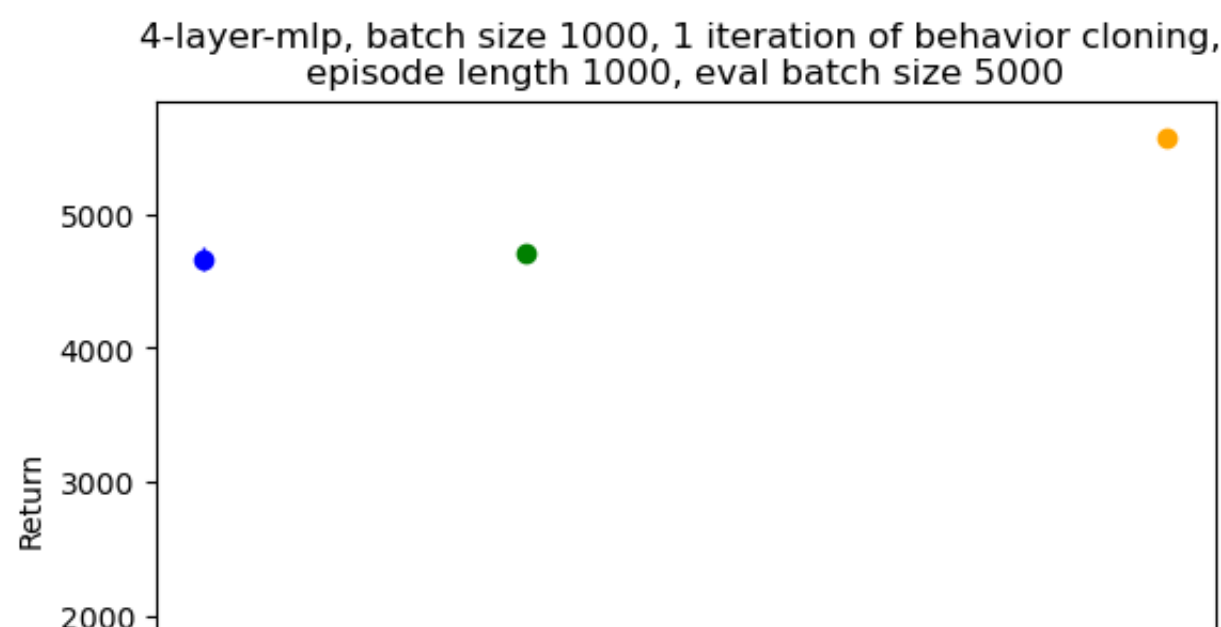
```
plt.errorbar(np.array(2), np.array(means[2]), np.array(std[2]), color='green', fmt='ok', lw=1)
```


/var/folders/80/mq_tqq2929b94txcm70_rbyr0000gn/T/ipykernel_93228/4274029093.py:33: UserWarning: color is redundantly defined by the 'color' keyword argument and the fmt string "ok" (-> color='k'). The keyword argument will take precedence.

```
plt.errorbar(np.array(3), np.array(means[3]), np.array(std[3]), color='red', fmt='ok', lw=1)
```

/var/folders/80/mq_tqq2929b94txcm70_rbyr0000gn/T/ipykernel_93228/4274029093.py:34: UserWarning: color is redundantly defined by the 'color' keyword argument and the fmt string "ok" (-> color='k'). The keyword argument will take precedence.

```
plt.errorbar(np.array(4), np.array(means[4]), np.array(std[4]), color='orange', fmt='ok', lw=1)
```





| | bc_ant | expert_ant | bc_walker2d | expert_walker2d |
|--------------------|------------------|------------------|-------------------|------------------|
| Average Return | 4669.265625 | 4713.6533203125 | 289.3726501464844 | 5566.845703125 |
| Standard Deviation | 81.988525390625 | 12.196533203125 | 86.68399810791016 | 9.237548828125 |
| Min Return | 4585.6728515625 | 4701.45654296875 | 6.183746337890625 | 5557.6083984375 |
| Max Return | 4798.67333984375 | 4725.849609375 | 440.492919921875 | 5576.08349609375 |

BC_Ant Achievement as percentage of Expert_Ant: 99%
BC_Walker2d Achievement as percentage of Expert_Walker2d: 5.2%

Plotting BA_Ant and BC_Walker2d Achievements

Q1.3

```
for i in 2000 3000 4000 5000 6000 7000 8000
do
python cs285/scripts/run_hw1.py
--expert_policy_file cs285/policies/experts/Walker2d.pkl
--env_name Walker2d-v4 --exp_name bc_walker2d --n_iter 1
--expert_data cs285/expert_data/expert_data_Walker2d-v4.pkl
--video_log_freq -1 --eval_batch_size 5000
--ep_len 1000 --num_agent_train_steps_per_iter $i
done
```

```
q1_bc_walker2d_Walker2d-v4_12-09-2022_00-51-27
q1_bc_walker2d_Walker2d-v4_12-09-2022_00-51-29
q1_bc_walker2d_Walker2d-v4_12-09-2022_00-51-32
q1_bc_walker2d_Walker2d-v4_12-09-2022_00-51-34
q1_bc_walker2d_Walker2d-v4_12-09-2022_00-51-37
q1_bc_walker2d_Walker2d-v4_12-09-2022_00-51-39
```

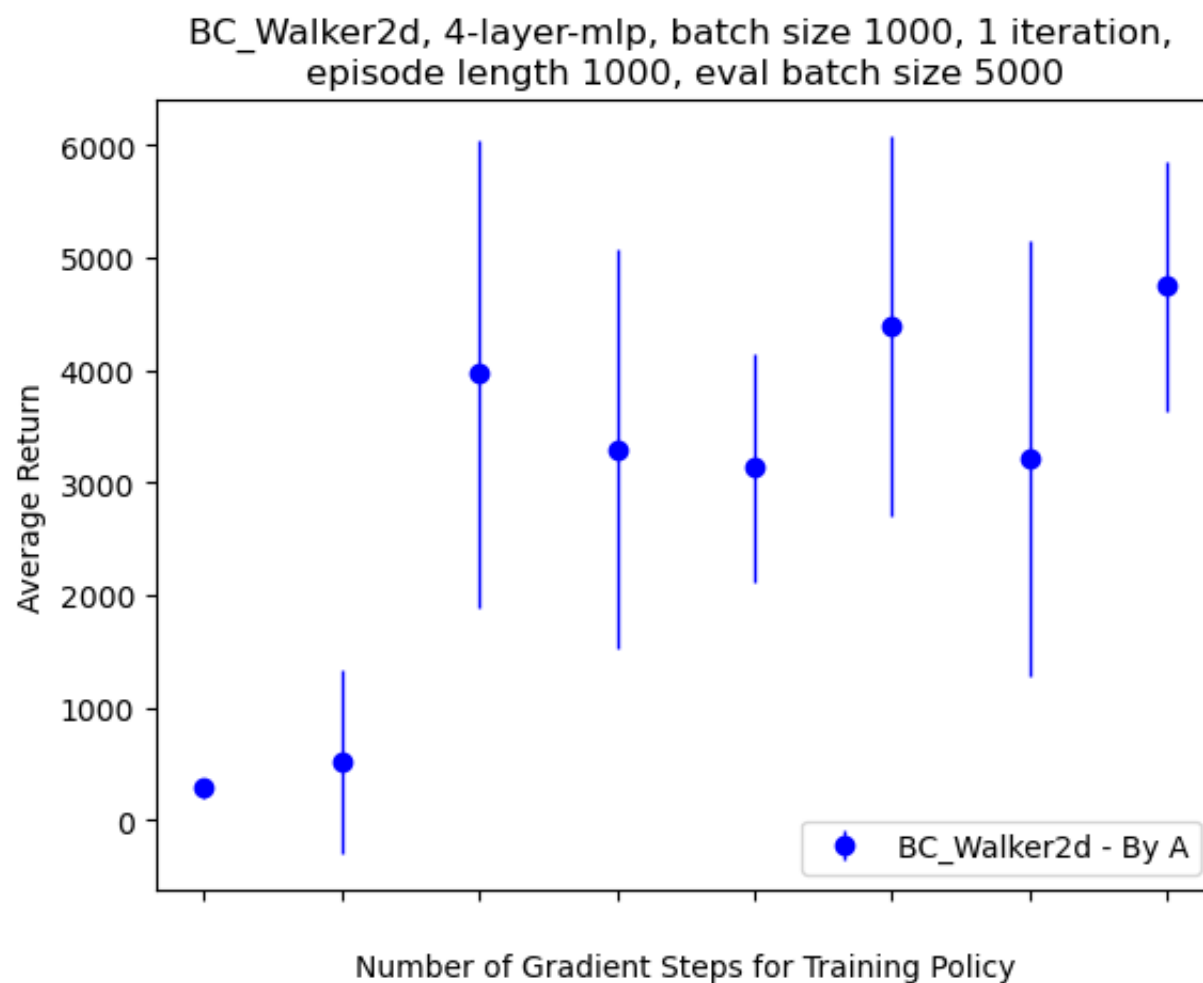
```

In [9]: 1 q_1_3_list_of_df = []
2 dir_name_prefix = 'q1_bc_walker2d_Walker2d-v4_12-09-2022_00-51-'
3 for i in [27,29,30,32,34,37,39]:
4     dir_name = dir_name_prefix + '{}'.format(i)
5     exp_df = to_df(dir_name)
6     q_1_3_list_of_df.append(exp_df)
7
8 means = [bc_walker2d_exp_Eval_AverageReturn]
9 stds = [bc_walker2d_exp_df["Eval_StdReturn"][0]]
10 for df in q_1_3_list_of_df:
11     Eval_AverageReturn = df["Eval_AverageReturn"][0]
12     Eval_StdReturn = df["Eval_StdReturn"][0]
13     means.append(Eval_AverageReturn)
14     stds.append(Eval_StdReturn)
15
16 plt.errorbar(np.array([1000, 2000, 3000, 4000, 5000, 6000, 7000, 8000 ]),
17             np.array(means), np.array(stds), color='blue', fmt='ok', lw=1)
18
19 plt.legend(['BC_Walker2d - By A'],loc='lower right')
20 plt.xlabel("Number of Gradient Steps for Training Policy")
21 plt.ylabel("Average Return")
22 plt.xticks(color='w')
23 plt.title("BC_Walker2d, 4-layer-mlp, batch size 1000, 1 iteration, \nepisode length 1000, eval batch
24 plt.show()
25

```

/var/folders/80/mq_tqq2929b94txcm70_rbyr0000gn/T/ipykernel_93228/3643757393.py:16: UserWarning: color is redundantly defined by the 'color' keyword argument and the fmt string "ok" (-> color='k'). The keyword argument will take precedence.

```
plt.errorbar(np.array([1000, 2000, 3000, 4000, 5000, 6000, 7000, 8000 ]),
```



Here I choose - Number of Gradient Steps for Training Policy - to be tuned. The MLP worked fine for the Ant environment leading me to believe that the model itself should be sufficient since the environments aren't too different.

The reason for imitation learning to not work well in the beginning could be due to divergence of training track too fast into the training. Thus, by providing more training data and more gradient steps to be performed on the agent we are likely to get closer trajectories to the training trajectory and hence better, longer, and more training down the line.

It seems the hyperparameter tuning is working well as average return is going up and eventually approached close performance to expert level.

Q2.2

q2_dagger_ant_Ant-v4_12-09-2022_01-29-46


q2_dagger_walker2d_Walker2d-v4_12-09-2022_01-31-24

```
python cs285/scripts/run_hw1.py
--expert_policy_file cs285/policies/experts/Ant.pkl
--env_name Ant-v4 --exp_name dagger_ant --n_iter 10
--do_dagger --expert_data cs285/expert_data/expert_data_Ant-v4.pkl
--video_log_freq -1 --eval_batch_size 5000 --ep_len 1000
```

```
python cs285/scripts/run_hw1.py
--expert_policy_file cs285/policies/experts/Walker2d.pkl
--env_name Walker2d-v4 --exp_name dagger_walker2d --n_iter 10
--do_dagger --expert_data cs285/expert_data/expert_data_Walker2d-v4.pkl
--video_log_freq -1 --eval_batch_size 5000 --ep_len 1000
```

```
In [10]: 1 dagger_ant_exp_path = "q2_dagger_ant_Ant-v4_12-09-2022_01-29-46"
2 dagger_ant_exp_df = to_df(dagger_ant_exp_path)
3
4 dagger_walker2d_exp_path = "q2_dagger_walker2d_Walker2d-v4_12-09-2022_01-31-24"
5 dagger_walker2d_exp_df = to_df(dagger_walker2d_exp_path)
```

```
In [11]: 1 dagger_ant_Eval_AverageReturn = dagger_ant_exp_df["Eval_AverageReturn"]
2 dagger_ant_Eval_StdReturn = dagger_ant_exp_df["Eval_StdReturn"]
3 dagger_ant_eval = dagger_ant_Eval_AverageReturn
4 expert_ant = np.ones(len(dagger_ant_Eval_AverageReturn)) * dagger_ant_exp_df["Train_AverageReturn"]
5 bc_ant = np.ones(len(dagger_ant_Eval_AverageReturn)) * bc_ant_exp_Eval_AverageReturn
6 plt.errorbar(np.array([0,1,2,3,4,5,6,7,8,9]), np.array(dagger_ant_Eval_AverageReturn),
7             np.array(dagger_ant_Eval_StdReturn), color='red', lw=1)
8
9 plt.plot(dagger_ant_eval, color = 'r')
10 plt.plot(expert_ant, color = 'orange')
11 plt.plot(bc_ant, color = 'magenta')
12 plt.ylim([4200,5000])
13 plt.xticks([0,1,2,3,4,5,6,7,8,9])
14 plt.legend(['dagger_ant', 'expert_ant', 'bc_ant'])
15 plt.xlabel("Dagger Iterations")
16 plt.ylabel("Return")
17 plt.title("Ant Env, 4-layer-mlp, batch size 1000, Dagger 10 iteration, \nepisode length 1000, eval b
18 plt.show()
```




```

In [12]: dagger_walker_Eval_AverageReturn = dagger_walker2d_exp_df["Eval_AverageReturn"]
dagger_walker_Eval_StdReturn = dagger_walker2d_exp_df["Eval_StdReturn"]
dagger_walker_eval = dagger_walker_Eval_AverageReturn
expert_walker = np.ones(len(dagger_walker_Eval_AverageReturn)) * dagger_walker2d_exp_df["Train_AverageReturn"]
bc_walker = np.ones(len(dagger_walker_Eval_AverageReturn)) * bc_walker2d_exp_Eval_AverageReturn
plt.errorbar(np.array([0,1,2,3,4,5,6,7,8,9]), np.array(dagger_walker_Eval_AverageReturn),
7         np.array(dagger_walker_Eval_StdReturn), color='red', lw=1)
8
9 plt.plot(dagger_walker_eval, color = 'r')
10 plt.plot(expert_walker, color = 'orange')
11 plt.plot(bc_walker, color = 'magenta')
12 #plt.ylim([3000,5000])
13 plt.xticks([0,1,2,3,4,5,6,7,8,9])
14 plt.legend(['dagger_walker', 'expert_walker', 'bc_walker'])
15 plt.xlabel("Dagger Iterations")
16 plt.ylabel("Return")
17 plt.title("Walker2d Env, 4-layer-mlp, batch size 1000, Dagger 10 iteration, \nepisode length 1000, eval batch size 5000")
18 plt.show()

```

